

A Comparison of k -Nearest Neighbours and Classification Trees for Multi-Class Network Intrusion Detection

G. Scott

Student number: 25935453

Email: 25935453@sun.ac.za

Abstract—Network intrusion detection is posed here as a multi-class classification problem, where each instance describes a single network connection labelled as normal traffic or as one of nine attack categories. The class distribution is heavily skew, and the rare attack categories are the ones an intrusion detector most needs to find. This report builds, tunes and compares a k -nearest neighbours classifier and a classification tree on such a dataset. Pre-processing was decided separately for each model, because a distance-based learner and a classification tree are not affected equally by the same data quality issue. Control parameters were tuned by grid search on a pool of instances held disjoint from the evaluation data, and both models were then evaluated by ten-fold stratified cross-validation on the remaining instances over identical folds. The classification tree recorded the higher macro-averaged F1, won every fold, and the difference survived a corrected resampled t -test. The ordering reversed on plain accuracy, so the better model depends on the value placed on the rare categories. Conflicting target levels on identical feature vectors also place a ceiling on macro-averaged F1. Neither model came close to that ceiling, so the gap between the two models is smaller than the gap between either model and the limit of the data.

I. INTRODUCTION

Network intrusion detection separates normal network traffic from attack traffic, and the task suits supervised classification. Labelled historical traffic is available for training, and new connections arrive continuously for classification. The dataset holds 257 673 instances, each describing a single network connection through 43 descriptive features. The target feature records normal traffic or one of nine attack categories, and the class distribution is heavily skew, with the largest level holding 534.5 times as many instances as the smallest. Several data quality issues also remain, including missing values, strong correlations, large positive outliers, numeric ranges spanning several orders of magnitude, a unique identifier and three nominal features.

The goal of this assignment is to build, tune and compare two classification models on the dataset. The first model is k -nearest neighbours (kNN), and the second is a classification tree. The two algorithms rest on different assumptions, so the same data quality issue does not affect them equally, and pre-processing was therefore decided separately for each model.

For both models the control parameters were tuned by grid search on a pool of 30 000 instances, held disjoint from the

data used for evaluation. The models were then evaluated by ten-fold stratified cross-validation on the remaining 227 673 instances. Performance was measured primarily by macro-averaged F1, which gives every target level equal weight. The tree reached 0.5749 against 0.4975 for kNN and won on every fold. The ordering reverses on plain accuracy, so the better model depends on the value placed on the rare attack categories. A third figure frames both scores: 15.7% of instances share a feature vector with an instance carrying a different label, so a macro-averaged F1 above 0.7759 is unreachable.

The remainder of the report is organised as follows: Section II describes the two classification algorithms and states the expectations regarding the remaining data quality issues. Section III describes the implementation and the pre-processing applied to each model. Section IV describes the empirical procedure. Section V reports and discusses the results, and Section VI concludes the report.

II. BACKGROUND

This section provides the background required for the remainder of the report. Section II-A describes the two classification algorithms at a high level. Section II-B states the expected impact of each remaining data quality issue on each of the two algorithms.

A. Similarity-Based and Information-Based Classification

Similarity-based learning predicts the target level of a query instance by finding the most similar instances in a labelled dataset [1]. The kNN algorithm is the standard example. Training consists only of storing the dataset, so kNN is a lazy learner [3]. A prediction is made by measuring the distance from the query to every stored instance, selecting the k closest, and returning the majority target level among them [1], [3].

Distance is measured using a metric from the Minkowski family, where a parameter value of one gives the Manhattan distance and a value of two gives the Euclidean distance [1], [3]. Euclidean distance squares each per-feature difference, so one large difference has a strong effect on the result. Manhattan distance sums the differences directly, so several smaller differences carry a more similar influence [3]. A plain majority vote gives all k neighbours equal weight. Distance-weighted

voting instead gives closer neighbours more influence through a decreasing function of distance [1].

Information-based learning builds a model by repeatedly partitioning the dataset on the descriptive feature that best separates the target levels [2]. A classification tree is the resulting structure, where each internal node tests one descriptive feature against one threshold and each leaf carries a predicted target level [3]. At every node the algorithm evaluates the candidate features and selects the split giving the largest reduction in impurity, then repeats the procedure on each resulting partition [2], [3].

Entropy and the Gini index both measure impurity, and both are lowest when a partition contains only one target level [3]. Neither measure wins in general, so comparing both on the dataset in use is recommended [3]. Growing a tree until every leaf is pure leads to overfitting, and pruning controls this by stopping growth early or by removing subtrees after the tree has been built [2], [3]. Wickramarachchi *et al.* described the partition produced by the classification and regression tree algorithm as axis-parallel, because each split tests a single feature, in contrast to oblique trees, where a split uses a linear combination of features [5].

kNN and the classification tree place their cost at opposite ends. kNN performs its distance calculations at prediction time, so the cost of a prediction grows with the size of the stored dataset [3]. A classification tree does most of its work during training, and a prediction only requires following a path from the root to a leaf.

B. Expectations Regarding Data Quality Issues

Eight data quality issues remain in the dataset. Their impact depends on the model's inductive bias, so expectations are considered separately for kNN and the classification tree. These expectations are made before any measurements, and Section V returns to them.

The identifier feature has a unique value for every instance. Information gain favours features with many possible outcomes, since splitting on them can produce many small, pure subsets [2], [3]. Wang and Witten describe the limiting case where an enumerated attribute has a different value for every instance and is selected first, even though it has no predictive power [6]. The tree is therefore expected to use the identifier as its root split and learn nothing useful. The identifier also has a much wider numeric range than most other features, so it would dominate distance calculations in kNN. Although the feature is useless to both models, the tree is expected to be affected more because it actively favours the identifier when selecting splits.

The numeric features span several orders of magnitude. Since distance is calculated over all descriptive features, a feature measured in bytes contributes much larger differences than one measured in seconds, leaving the smaller-scale features with very little influence [1], [3]. The classification tree is expected to be unaffected, because thresholds are selected by sorting the instances on one feature and testing boundaries between adjacent instances with different target levels [3]. A

strictly increasing rescaling does not change that ordering, so the partition stays the same. Scaling is therefore expected to matter for kNN alone.

Several features are perfectly correlated. The specification identifies sbytes with sloss, dbytes with dloss, and is_ftp_login with ct_ftp_cmd, and further strong correlations are present. A feature that strongly correlates with another is redundant, because no additional information is carried [3]. A duplicated feature is therefore counted twice in the distance and carries double weight. kNN is also sensitive to the curse of dimensionality, whereas a tree selects a subset of the features during induction [3]. Removing redundant features is therefore expected to improve kNN and to have little effect on the classification tree.

Many features carry large positive outliers. For kNN, an outlier lies far from the query by construction, so an outlier is unlikely to enter a small neighbourhood, and a majority vote limits the influence of any outlier that does [1]. For the tree, an outlier is isolated in a small leaf during growth. Pruning then removes the leaf and absorbs the instance into a parent partition where it sits in the minority [2], [3]. Both models are expected to tolerate outliers, but the argument for the tree holds only if the tree is pruned.

The service feature is missing for 54.8% of instances. A distance is still calculable from the features that are present, but the approach is only reliable when few values are missing [1], so more than half is expected to be a problem. Tree induction handles missing values directly, either by distributing an instance across outcomes using fractional weights [2] or by splitting on a correlated feature [6], so the tree is expected to cope better. Since no instance is missing more than one value, the missing service values are likely systematic rather than random. An absent service is most plausibly a connection with no identifiable application-layer service, making it an unobserved category rather than a lost value.

The class distribution is heavily skewed. For the tree, minority classes produce small leaves that pruning can remove, causing those instances to be absorbed into a majority partition [2]. For kNN, a large value of k can also cause the majority class to dominate the vote [3]. Distance weighting reduces this effect but does not remove it, since a sufficiently imbalanced dataset can still favour the majority class [3]. The Worms class is the most extreme example, holding less than 0.1% of instances. A Worms query is unlikely to hold a Worms instance among its nearest neighbours. Both models therefore need an explicit correction, and kNN is expected to need it more.

Three features are nominal, where the proto feature carries 133 levels, against 13 for service and 9 for state, so proto dominates whatever encoding is chosen. Categorical similarity is measured by overlap or by an index such as Jaccard [1], [4]. One-hot encoding produces sparse binary features, in which most pairwise comparisons are co-absences, and the recommended index in that situation is one that ignores co-absences [3]. For the tree, one-hot encoding splits the impurity reduction available from proto across 133 separate

binary features, so no single feature looks worth splitting on. An ordinal encoding keeps the levels in one feature and is therefore expected to suit the tree better.

Finally, identical feature vectors carry conflicting target levels. No classifier is able to separate two instances that agree on every descriptive feature, so the issue places a limit on what any model reaches.

In summary, the identifier, the differing numeric ranges, the correlations and the nominal encoding are expected to affect kNN most, and the identifier and the skew class distribution are expected to affect the tree most.

III. METHODOLOGY

This section describes how the two models were implemented and how the data was prepared for each. Section III-A describes the implementation and the libraries used. Section III-B presents the pre-processing decisions taken for each model, together with the justification for each decision.

A. Implementation

Both models were implemented in Python using scikit-learn¹ version 1.7.2, with pandas² version 2.3.0 for data handling and NumPy³ version 2.2.6 for numeric work. kNN uses the KNeighborsClassifier class, and the tree uses the DecisionTreeClassifier class, an optimised implementation of the CART algorithm.

Each model was implemented as a two-stage pipeline. The first stage is a column transformer that applies one set of pre-processing steps to the numeric features and another to the nominal features, and the second stage is the classifier. Treating the pipeline as a single estimator ensures that the quantile mapping and the encoder levels are fitted on the training fold alone.

The cross-validation loop was implemented directly instead of through a library helper, so that both pipelines run over one set of fold indices and the per-fold scores, per-class scores, timings and confusion counts are collected in one pass.

The code is separated by phase, covering exploration, pre-processing, tuning, evaluation, and two diagnostic scripts that compute the ambiguity ceiling and the average class accuracy. The source is available in the repository⁴.

B. Data Pre-Processing

Pre-processing followed one principle: a transformation is applied only where the algorithm receiving the data requires it. Leaving an issue untreated is as much a decision as transforming it, and every such decision is justified below. Table ?? lists the decision taken for each issue under each model, and the reasoning behind each row follows in the text.

1) *Shared Cleaning*: Four cleaning rules were applied before either pipeline, but only one changed the data. The identifier was removed, reducing the dataset from 44 columns to 43 and leaving 42 descriptive features. The other three rules removed nothing, since there were no missing target values, no constant features, and no instance with more than one missing value. The issues covered by those rules had been introduced into the dataset of assignment one and have since been removed.

2) *Transformations Applied*: One-hot encoding of proto, state and service produces 155 binary features, with 133 coming from proto alone. Together with the 39 numeric features, this gives 194 columns in the evaluation pool. The tuning pool has 192, because two levels only occur in the larger pool. The encoders are fitted separately on each training fold, with unknown levels ignored. This prevents levels that are missing from a training fold from causing errors when transforming its test fold. Both models therefore use the same feature space, as the tuned tree also uses one-hot encoding.

The numeric features are not scaled for the tree, for the reason given in Section II-B: a monotone rescaling cannot change a partition built from orderings alone.

No outliers were removed from either model. The task is to detect attacks, and attack traffic is precisely what produces the large positive values, so removing those values would discard the signal the models are meant to find. For the tree the pruning parameters already dampen the influence of a rare extreme value [2], [3].

No correlation filter was applied, against the expectation of Section II-B. Kelleher *et al.* distinguished filters from wrappers: a filter evaluates features individually, so redundant features are kept or useful combinations removed, whereas a wrapper evaluates feature subsets by the performance of a model trained on them, taking the inductive bias of that model into account [3]. The filter was therefore tested as a wrapper across thresholds of 0.90, 0.95 and 0.99 with three scalers, and the unfiltered pipeline came top within every scaler family. Under the selected quantile transformation, macro F1 fell from 0.4465 without filtering to 0.4437, 0.4376 and 0.4373 at the three thresholds. No gap exceeded the fold variation, so the filter is reported as never helping rather than as harmful. A likely reason is that the filter drops whole features, which corrects the doubled weight of pairs such as sbytes and sloss but also removes features that still carry useful information.

Median imputation appears in both pipelines but never changes a value, because service is the only feature with missing values and service is nominal. The step exists only because scikit-learn rejects missing values, which means the native handling described in Section II-B was unavailable, since the estimator supports neither fractional instance weights [2] nor surrogate splitting [6].

No resampling was applied. Balanced class weighting reweights the minority classes without removing majority instances or creating additional minority samples, and resampling could place duplicate instances in both training and test folds, which would inflate the cross-validation results.

¹<https://scikit-learn.org>

²<https://pandas.pydata.org>

³<https://numpy.org>

⁴https://github.com/Gabriella-Scott/ML_Assignments/tree/main/Assignment2

In conclusion, kNN receives scaling and encoding, the tree receives encoding alone, and both are otherwise left untouched. Three of the transformations considered were rejected on measurement rather than on argument.

IV. EMPIRICAL PROCEDURE

This section describes the procedure followed to obtain the reported results. Section IV-A describes the control parameter tuning. Section IV-B motivates the performance metrics. Section IV-C describes the evaluation and the statistical comparison.

A. Control Parameter Tuning

The usable instances were divided into a stratified tuning pool of 30 000 and a separate evaluation pool of 227 673. This ensured that no instance used to select a parameter was later used to evaluate the model. Testing on data used during training can give an overly optimistic performance estimate [3]. This is especially important here because the correlation filter was chosen using a wrapper. Since a wrapper trains models to evaluate its choices, it requires its own data, while a separate evaluation set must still be kept aside [3].

Both models were tuned using grid search with five-fold stratified cross-validation on the tuning pool, using macro F1 as the selection metric. The search was split into two stages. The first selected the data representation, including the scaler and correlation threshold for kNN, and the encoding, split criterion and class weighting for the tree. The second stage tuned the remaining classifier parameters while holding the selected representation fixed. Searching all combinations at once would have required 432 configurations for kNN and 1 120 for the tree. Staging the search reduced that cost, although the approach is greedy and may miss interactions between the two groups of parameters.

A small k overfits noise while a large one underfits, and no general rule fixes the value [3]. Entropy and the Gini index were both searched because the more effective impurity measure varies from one dataset to another [3]. Depth and minimum leaf size stunt growth early, whereas the cost complexity parameter prunes the finished tree, so the grid covers both routes to controlling overfitting [3]. Distance weighting follows the Shepard scheme [1], and the distance metric was searched rather than fixed in advance, since a similarity measure that wins on one task need not win on another [4].

Table I gives the selected configuration for each model. Not every difference in that table is real. Manhattan beat Euclidean at all nine values of k under both weighting schemes, so that selection holds. Others sit inside the noise. A value of three for k scored 0.4671 ± 0.0147 against 0.4644 ± 0.0052 for a value of one, and a depth of 15 scored 0.5324 ± 0.0266 against 0.5322 for a depth of 20. Both are reported as selected rather than as better. Depth 15 remains the sounder choice, because its training macro F1 was 0.691 against 0.774 at depth 20 for the same validation score, making it the simpler model.

Figure ?? shows how macro F1 responds to the main control parameters. The kNN curve falls steadily beyond k of five,

TABLE I
CONTROL PARAMETERS SEARCHED AND SELECTED

Parameter	Values searched	Selected
<i>k-nearest neighbours, macro F1 0.4671 ± 0.0147</i>		
Scaler	standard, robust, quantile	quantile
Filter	none, 0.99, 0.95, 0.90	none
k	1, 3, 5, 7, 9, 11, 15, 21, 31	3
Vote weight	uniform, distance	distance
Metric	Euclidean, Manhattan	Manhattan
<i>Classification tree, macro F1 0.5324 ± 0.0266</i>		
Encoding	ordinal, one-hot	one-hot
Criterion	Gini, entropy	entropy
Class weight	none, balanced	balanced
Maximum depth	5, 10, 15, 20, 25, 30, none	15
Minimum leaf	1, 2, 5, 10, 20	2
Pruning	0, 1e-5, 1e-4, 1e-3	1e-4

which is the expectation of Section II-B appearing directly in the tuning data: a wider neighbourhood lets the majority class take the vote. The tree shows the opposite failure at the other end of its range. At a depth of 5 the training and validation curves sit together at roughly 0.32, an unambiguous underfit, and past a depth of 15 the training curve climbs steadily while the validation curve stays flat. Cost complexity pruning made no measurable difference, with all four values tested falling inside each other's error bars.

B. Performance Metrics

Accuracy was not used as the primary metric because it can hide poor performance when the target classes are imbalanced [3]. The tuning results show this clearly. With the encoding and split criterion fixed, changing class weighting from none to balanced reduced accuracy from 0.7771 to 0.7691. However, macro F1 increased from 0.4898 to 0.5143, while balanced accuracy increased from 0.4939 to 0.5255. The improvement in minority-class performance was therefore barely reflected in accuracy.

Precision measures the proportion of predictions for a class that are correct, while recall measures the proportion of instances from that class that are correctly identified. F1 is the harmonic mean of precision and recall, so a low value in either measure has a stronger effect on the result [3]. For the 10 target classes, these metrics are calculated separately for each class using the confusion matrix [3].

Macro F1, the unweighted mean of the per-class F1 scores, is therefore the primary metric, since Worms carries the same weight as Normal despite the difference in size between them. Average class accuracy on a harmonic mean is recommended as the default for a categorical target [3], so it is reported alongside. Average class accuracy aggregates recall alone, whereas macro F1 aggregates precision and recall together, and the distinction matters here because balanced class weighting buys recall by spending precision.

Query time is reported as well, since how quickly a model produces a prediction is part of how useful it is [3].

C. Analysis Procedure

Both models were evaluated using ten-fold stratified cross-validation on the evaluation pool. A single hold-out split can give a misleading result because the estimate depends on which instances happen to be selected for testing. Cross-validation reduces this effect by averaging performance across several splits [3]. Ten folds is also a common choice in practice [3]. Stratification keeps the class proportions similar across folds, which is important when the smallest class has only around 15 instances per fold.

The comparison is paired. One StratifiedKFold object produced the indices and both pipelines ran over the same ten splits, so each fold yields two scores measured on identical instances, and the ten differences form the sample for the test. A plain paired t -test would overstate the outcome, because the training sets of the folds overlap heavily and the fold scores are therefore not independent. The corrected resampled test adds the test-to-train size ratio to the variance term, giving $1/k + 1/(k - 1)$ for k folds⁵. Both statistics are reported so that the effect of the correction is visible.

The ambiguity ceiling was identified during a diagnostic check rather than planned in advance. A one-nearest-neighbour model evaluated on its own training data should achieve perfect accuracy because each instance is its own closest neighbour [3]. Instead, the model scored 0.853, which indicates that some identical feature vectors have different labels. Instances were then grouped by their exact feature vectors. If a group contains multiple labels, the best possible prediction is the majority label for that group. Using this label for every instance gives an upper bound on performance. The full dataset contains 153,684 distinct vectors, with 15.7% of instances belonging to groups with conflicting labels. This gives an upper bound of 0.9139 accuracy and 0.7759 macro F1. Since only exact duplicate vectors are considered, the true maximum performance may be lower.

V. RESEARCH RESULTS

This section reports and discusses the results obtained under the procedure of Section IV. Section V-A reports aggregate performance and the paired comparison, Section V-B examines the per-class results against the achievable ceiling, and Section V-C returns to the expectations of Section II-B.

A. Aggregate Performance

Table II reports the mean and standard deviation of each metric over the ten folds. Average class accuracy is computed from the confusion matrix summed across folds [3], so it carries no standard deviation.

On the primary metric the tree won clearly. It reached a macro F1 of 0.5749 against 0.4975, and won on all ten folds individually. The mean difference was 0.0774 with a standard deviation of 0.0108, giving an uncorrected t statistic of 22.55 and a corrected statistic of 15.52. Both fall well beyond

TABLE II
AGGREGATE PERFORMANCE OVER TEN FOLDS

Metric	Classification tree	k -nearest neighbours
Macro F1	0.5749	0.4975
Standard deviation	0.0061	0.0089
Avg. class accuracy (AM)	0.6841	0.4912
Avg. class accuracy (HM)	0.5711	0.2017
Macro precision	0.5540	0.5389
Macro recall	0.6841	0.4912
Accuracy	0.7481	0.7917
Weighted F1	0.7759	0.7908
Fit time per fold (s)	11.5	2.3
Predict time per fold (s)	0.07	316.5

any conventional significance threshold, so the conclusion survives the more conservative test rather than depending on the optimistic one.

Which model performs better depends on the metric. kNN achieved higher accuracy, at 0.7917 against 0.7481, and higher weighted F1, at 0.7908 against 0.7759, while the tree led on macro F1 and on both measures of average class accuracy. Figure 1 shows why. Balanced class weighting makes the tree more willing to sacrifice majority-class instances to improve minority recall, and it misclassified 27.4% of Normal traffic, with 22.3% sent to Fuzzers. Normal traffic makes up 36% of the evaluation pool, so that misrouting accounts for much of the accuracy gap.

The difference becomes much larger when using the harmonic mean. The tree’s macro F1 is 1.16 times higher than kNN’s, while its harmonic mean of class accuracy is 2.8 times higher, at 0.5711 compared with 0.2017. The arithmetic mean gives a smaller gap of 0.6841 compared with 0.4912. Since the harmonic mean is strongly affected by low values, kNN’s recall of only 5.2% on Worms and 10.6% on Backdoor pulls its score down substantially. In contrast, the arithmetic mean allows high recalls such as 92.1% on Normal and 97.7% on Generic to hide these weaker classes. The recommended metric therefore also provides the clearest separation between the two models.

B. Per-Class Performance

Table III gives the per-class F1 of both models against the ceiling imposed by the conflicting labels. Reading the weak classes as a single group would be a mistake, because they are weak for two different reasons, and the two call for different remedies.

Backdoor, Analysis and DoS are limited by the data itself, since their ceilings are 0.261, 0.375 and 0.528, so no classifier can achieve a high score on these classes. Measured against what is actually achievable, the tree performs reasonably well on DoS, with a score of 0.464, which is 88% of the available 0.528. The confusion matrix shows that these three classes are often confused with each other, with 51.0% of Backdoor instances and 49.7% of Analysis instances predicted as DoS. This is mainly due to the features available. Every descriptive feature is a connection-level statistic, covering duration, byte

⁵https://scikit-learn.org/stable/auto_examples/model_selection/plot_grid_search_stats.html

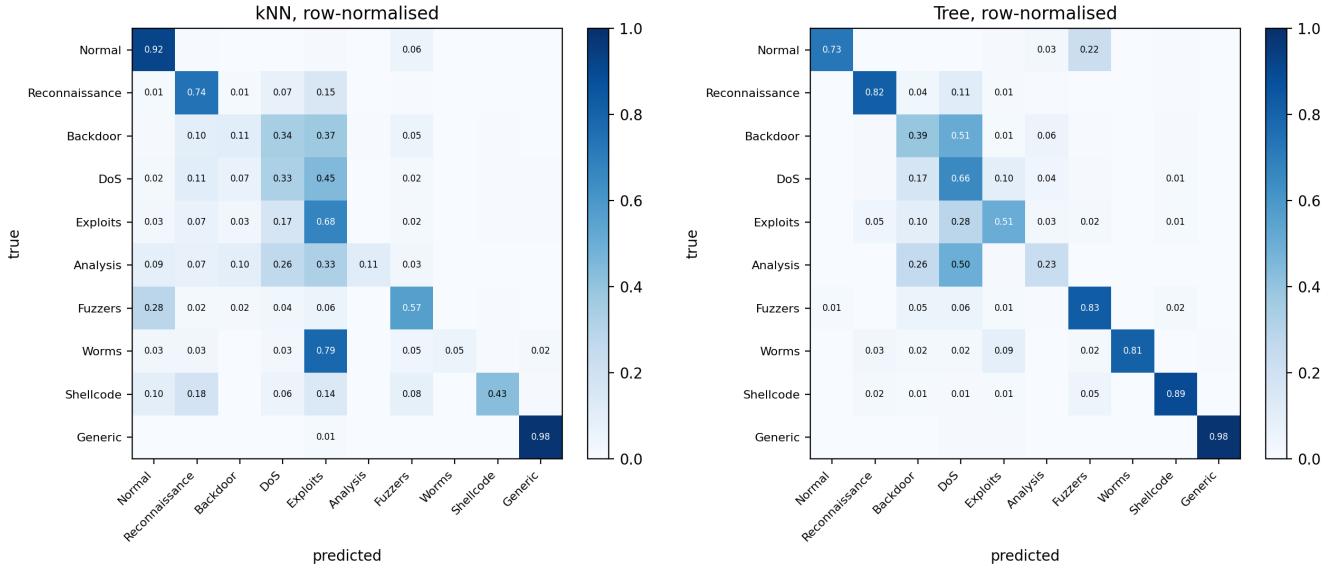


Fig. 1. Confusion matrices for the classification tree and for the nearest neighbour model

TABLE III
PER-CLASS F1 AGAINST THE ACHIEVABLE CEILING

Class	Classification tree	k -nearest neighbours	Ceiling
Normal	0.839	0.913	0.997
Generic	0.987	0.985	0.997
Reconnaissance	0.823	0.672	0.913
Exploits	0.637	0.675	0.817
Fuzzers	0.607	0.606	0.914
Worms	0.576	0.067	0.954
Shellcode	0.528	0.463	0.998
DoS	0.464	0.328	0.528
Analysis	0.149	0.188	0.375
Backdoor	0.139	0.079	0.261

and packet counts, timing and protocol flags. None of them describes what the attacker was trying to do. If two attacks have different intentions but produce similar connection patterns, the models cannot distinguish them because that information is not present in the features.

Shellcode and Worms have the opposite problem. Their ceilings are 0.998 and 0.954, so both are almost perfectly separable, but the evaluation pool holds only 1335 and 154 instances of them. The weakness is representation rather than separability, so more examples would help, whereas additional data would not solve the feature limitation affecting Backdoor and Analysis.

Set against the ceiling, the tree reached 74% of the achievable macro F1 and kNN 64%. The distance between the two models is therefore smaller than the distance between either model and the limit of the data.

C. Confirmation of the Stated Expectations

The prediction about Worms in Section II-B held. kNN scored an F1 of 0.067 on the class, recovering 5.2% of its instances and sending 78.6% of them to Exploits. With a

class that small a Worms query rarely holds even one Worms instance among its three nearest neighbours, and KNeighborsClassifier offers no class weighting to compensate, which is the predicted outcome of a sufficiently imbalanced dataset letting the majority level dominate despite distance weighting [3]. The tree, which could be weighted, reached 0.576 with a recall of 81.2%.

Weighting was not free. The tree bought its recall by spending precision, ending with a macro recall of 0.684 against a macro precision of 0.554. Backdoor shows the trade at its most extreme: the tree predicted the class 9551 times for 2058 actual instances, roughly four and a half times too often, for a precision of 0.084 against a recall of 0.392. A model tuned on macro F1 accepts that exchange because finding a rare attack usually matters more than occasionally raising a false alarm.

The expectation about nominal encoding was not supported. One-hot encoding reached 0.5143 against 0.4991 for ordinal encoding under the selected criterion and weighting, a margin of roughly one standard deviation, and one-hot won three of the four criterion and weighting combinations tested. The reasoning in Section II-B overlooked how the split is actually found. The best split for an enumerated attribute with k values lies at one of the $k-1$ positions obtained by ordering the levels by their average class value [6], but OrdinalEncoder orders levels alphabetically. Across 133 protocol levels the tree was therefore searching 132 cut points in an essentially arbitrary ordering, which is worse than searching 133 binary features.

Query cost behaved as the lazy learning argument of Section II-A predicts. kNN needed 316.5 seconds per fold against 0.07 seconds for the tree, a difference of more than three orders of magnitude [3].

VI. CONCLUSIONS

This report compared k -nearest neighbours (kNN) with a classification tree on a ten-class network intrusion detection problem. Pre-processing was selected separately for each model, with control parameters tuned using grid search on a separate pool of 30 000 instances. Both models were then evaluated using ten-fold stratified cross-validation on the remaining 227 673 instances.

Overall, the classification tree was the better model. It achieved a higher macro F1 of 0.5749 compared with 0.4975, performed better on average class accuracy using both the arithmetic and harmonic means, and achieved higher F1 scores for seven of the ten classes. It was also more than three orders of magnitude faster at prediction time. The tree won all ten folds, and the difference remained significant under the corrected resampled t -test, not just the uncorrected version.

There is one important qualification. kNN achieved higher plain accuracy and weighted F1, but both metrics are dominated by the two largest classes, whereas macro F1 gives the rare attack classes equal importance. For intrusion detection that is the more useful priority, but it remains a choice of evaluation priority rather than an objective property of the dataset.

Two expectations drawn from theory were not supported. Removing correlated features was expected to improve kNN, and ordinal encoding was expected to suit the tree better than one-hot encoding. Neither held, and both were caught because the choices were tested during tuning.

The binding constraint is the data itself. The tree reached only 74% of the macro F1 ceiling imposed by the conflicting labels, so further work should begin with features that separate Backdoor, Analysis and DoS, which connection-level statistics alone cannot distinguish. A similarity index that ignores co-absences would suit the sparse binary features produced by one-hot encoding [3], and a k -d tree index would cut the query cost if kNN were ever deployed [3]. Any deployed model would also need continued validation, since attack traffic changes over time and models are subject to concept drift [3].

REFERENCES

- [1] A. P. Engelbrecht, “Similarity-based learning,” Topic 3 lecture notes, Machine Learning 741, Stellenbosch University, 2026.
- [2] A. P. Engelbrecht, “Information-based learning,” Topic 4 lecture notes, Machine Learning 741, Stellenbosch University, 2026.
- [3] J. D. Kelleher, B. Mac Namee, and A. D’Arcy, *Fundamentals of Machine Learning for Predictive Data Analytics: Algorithms, Worked Examples, and Case Studies*. Cambridge, MA: MIT Press, 2015, pp. 162–174, 182, 184–188, 196–199, 221, 230–233, 241–245, 251–254, 263–269, 426–427, 432–438, 442–445, 464, 472.
- [4] K. Rodden, W. Basalaj, D. Sinclair, and K. Wood, “A comparison of measures for visualising image similarity,” in *Proceedings of the Challenge of Image Retrieval*, 2000.
- [5] D. C. Wickramarachchi, B. L. Robertson, M. Reale, C. J. Price, and J. Brown, “HH CART: an oblique decision tree,” *Computational Statistics and Data Analysis*, vol. 96, pp. 12–23, 2016.
- [6] Y. Wang and I. H. Witten, “Inducing model trees for continuous classes,” in *Proceedings of the Poster Papers of the European Conference on Machine Learning*, 1997.